

# Transfer Learning on Image Data

The purpose of this set of notes is to give you background information for your next programming assignment.

## 1 Assignment

Create a program that classifies images of bees versus ants. You should follow the PyTorch tutorial on “Transfer Learning for Computer Vision” located at

[https://pytorch.org/tutorials/beginner/transfer\\_learning\\_tutorial.html](https://pytorch.org/tutorials/beginner/transfer_learning_tutorial.html)

I strongly recommend working through the “Introduction to PyTorch” tutorials before completing this tutorial. Start with the “Tensors” tutorial and work your way through the “Save and Load a Model” tutorial.

At the end of the transfer learning tutorial, you will generate a figure demonstrating the output of your code. You should upload your completed python file and this image to sakai.

The due date is next Sunday (20 Nov).

## 2 Image Data

Bees vs Ants:

- 2 classes
- 398 images
- Pre-determined train/validation split (245/153 images)

ImageNet dataset:

- 1000 classes, 1.2 million images
- test set size = 150,000

Image data is usually “noisy”

- Non-statistical “noise”
  - images are all of different resolutions
  - size of the bug is different for each image
  - some images have effects applied to them
- Statistical “noise”
  - scraped from the web, so there’s errors
  - human annotations, and humans make mistakes

Karpathy: human error rate is about 5%

<http://karpathy.github.io/2014/09/02/what-i-learned-from-competing-against-a-convnet-on-imagenet/>

- some images can legitimately have multiple labels


Reference: <https://www.unite.ai/assessing-the-historical-accuracy-of-imagenet/>

### 3 Deep Learning

Reference: <https://cs231n.github.io/>

- All famous deep learning image classification models are all trained on ImageNet
- ImageNet performance change over time:  
<https://paperswithcode.com/sota/image-classification-on-imagenet>
- Best results use ensembles of deep learning algorithms  
For example: “Deep Residual Learning for Image Recognition”

What deep learning looks like:



See also ResNet: “Deep Residual Learning for Image Recognition”


## Theoretical Problems

- Optimization: Highly non-convex optimization problems

- Statistical: VERY high VC dimensions

VC dimensions are VERY difficult to directly compute, so people typically report the number of parameters used in the model as a “proxy” for the VC dimension.

Double Descent Reference: <https://medium.com/mllearning-ai/double-descent-8f92dfdc442f>



## 4 Tricks for our dataset

Three tricks:

- 1. Use an ImageNet-trained CNN for feature extraction
  - “transfer knowledge” from the ImageNet task to the bees vs ants task
  - greatly reduces the VC dimension
- Computational notes:
  - pytorch’s **CrossEntropyLoss** is the logistic loss
  - Lots of variations of SGD (Momentum, Adam, AdaGrad)
  - Batch size
  - Learning rate scheduler
  - Still expensive computationally

## 2. Solution 2: dataset augmentation

- Reference: [http://d2l.ai/chapter\\_computer-vision/image-augmentation.html](http://d2l.ai/chapter_computer-vision/image-augmentation.html)
- “effective number of data points” equals number of epochs times training set size
- technically these new data points are not independent, and so we’re not “really” increasing our dataset in the VC dimension sense... but it’s really close and works well in practice

## 3. Solution 3: add regularization

- weight decay
- early stopping
- dropout
- residual layers
- batch normalization