

Notes: Feature Engineering

General references:

1. <https://www.kaggle.com/notebooks?sortBy=voteCount&group=everyone&pageSize=20&datasourceType=competitions>
2. <https://www.kaggle.com/shivamb/extensive-text-data-feature-engineering>
3. <https://www.kaggle.com/sudalairajkumar/getting-started-with-text-preprocessing>
4. <https://towardsdatascience.com/understanding-feature-engineering-part-3-traditional-methods-for-text>
5. Textbook, Section 25.2

Important: Keep in mind the differences between bayes error, approximation error, and estimation error throughout this discussion.

1 Text

Problem 1. What is the difference between a *distributed* encoding and a *1-hot* encoding?

Problem 2. A 1-hot encoding of words is sometimes called a *bag of words*. What are its limitations?

1. Large dimensionality

2. The context problem (i.e., not bijective)

3. The phrase problem

4. The synonym problem

5. The conjugation problem

6. The tokenization problem

7. The punctuation problem

See: <https://digitalsynopsis.com/tools/punctuation-marks-importance-rules-usage/>

8. The Unicode problem

See the video “Unicode and Python: the absolute minimum you need to know”: <https://www.youtube.com/watch?v=oXVmZGN6p1Y>

Problem 3. What are n -grams? What are the tradeoffs of using n -grams?

Problem 4. What is lemmatization? What are the tradeoffs of using lemmatization?

Problem 5. What is text normalization? What are the tradeoffs?

Problem 6. What is stop word elimination? What are the trade-offs of stop word elimination?

Problem 7. What is the TF-IDF transform? What are the tradeoffs?

Problem 8. What is the hashing trick? What are the trade-offs of using the hashing trick?

References:

1. Hashing trick tutorial: <https://booking.ai/dont-be-tricked-by-the-hashing-trick-192a6aae3087>
2. Zipf's law: https://en.wikipedia.org/wiki/Zipf%27s_law
3. Excellent research paper on the Johnson-Lindenstrauss lemma: <https://papers.nips.cc/paper/7784-fully-understanding-the-hashing-trick>

2 Time

Problem 9. The discretization transform.

Problem 10. The sin/cos transform.

3 Graph metadata

Problem 11. Friendship features.

Problem 12. How can pagerank be used in twitter classification?

4 Generic

Problem 13. What is the unit-normalization transform? What are the tradeoffs?

Problem 14. What is the clipping transform? What are the tradeoffs?

Problem 15. What is the log transform? What are the tradeoffs?

Problem 16. What is the whitening transform? What are the tradeoffs?